

Homework 10 - CIFAR10 Image Classification with PyTorch

About

The goal of the homework is to train a convolutional neural network on the standard CIFAR10 image classification dataset.

When solving machine learning tasks using neural networks, one typically starts with a simple network architecture and then improves the network by adding new layers, retraining, adjusting parameters, retraining, etc. We attempt to illustrate this process below with several architecture improvements.

▼ Dev Environment

Working on Google Colab

You may choose to work locally or on Google Colaboratory. You have access to free compute through this service. Colab is recommended since it will be setup correctly and will have access to GPU resources.

1. Visit <https://colab.research.google.com/drive>
2. Navigate to the **Upload** tab, and upload your HW10.ipynb
3. Now on the top right corner, under the **Comment** and **Share** options, you should see a **Connect** option. Once you are connected, you will have access to a VM with 12GB RAM, 50 GB disk space and a single GPU. The dropdown menu will allow you to connect to a local runtime as well.

Notes:

- If you do not have a working setup for Python 3, this is your best bet. It will also save you from heavy installations like `tensorflow` if you don't want to deal with those.
- *There is a downside.* You can only use this instance for a single 12-hour stretch, after which your data will be deleted, and you would have to redownload all your datasets, any libraries not already on the VM, and regenerate your logs.

Installing PyTorch and Dependencies

The instructions for installing and setting up PyTorch can be found at <https://pytorch.org/get-started/locally/>. Make sure you follow the instructions for your machine. For any of the remaining libraries used in this assignment:

- We have provided a `hw8_requirements.txt` file on the homework web page.
- Download this file, and in the same directory you can run `pip3 install -r hw8_requirements.txt`. Check that PyTorch installed correctly by running the following:

```
import torch
torch.rand(5, 3)

tensor([[0.6529, 0.8300, 0.3982],
        [0.7574, 0.3594, 0.1487],
        [0.6654, 0.7324, 0.2646],
        [0.7374, 0.7814, 0.9135],
        [0.2481, 0.2306, 0.9768]])
```

▼ Part 0 Imports and Basic Setup (5 Points)

First, import the required libraries as follows. The libraries we will use will be the same as those in HW8.

```
import numpy as np
import torch
from torch import nn
from torch import optim

import matplotlib.pyplot as plt
```

GPU Support

Training of large network can take a long time. PyTorch supports GPU with just a small amount of effort.

When creating our networks, we will call `net.to(device)` to tell the network to train on the GPU, if one is available. Note, if the network utilizes the GPU, it is important that any tensors we use with it (such as the data) also reside on the GPU. Thus, a call like `images = images.to(device)` is necessary with any data we want to use with the GPU.

Note: If you can't get access to a GPU, don't worry to much. Since we use very small networks, the difference between CPU and GPU isn't large and in some cases GPU will actually be slower.

```
import torch.cuda as cuda

# Use a GPU, i.e. cuda:0 device if it available.
device = torch.device("cuda:0" if cuda.is_available() else "cpu")
print(device)
```

cuda:0

▼ Training Code

```
import time

class Flatten(nn.Module):
    """NN Module that flattens the incoming tensor."""
    def forward(self, input):
        return input.view(input.size(0), -1)

def train(model, train_loader, test_loader, loss_func, opt, num_epochs=10):
    all_training_loss = np.zeros((0,2))
    all_training_acc = np.zeros((0,2))
    all_test_loss = np.zeros((0,2))
    all_test_acc = np.zeros((0,2))

    training_step = 0
    training_loss, training_acc = 2.0, 0.0
    print_every = 1000

    start = time.clock()

    for i in range(num_epochs):
        epoch_start = time.clock()

        model.train()
        for images, labels in train_loader:
            images, labels = images.to(device), labels.to(device)
            opt.zero_grad()

            preds = model(images)
            loss = loss_func(preds, labels)
            loss.backward()
```

```

opt.step()

training_loss += loss.item()
training_acc += (torch.argmax(preds, dim=1)==labels).float().mean()

if training_step % print_every == 0:
    training_loss /= print_every
    training_acc /= print_every

    all_training_loss = np.concatenate((all_training_loss, [[training_step,
    all_training_acc = np.concatenate((all_training_acc, [[training_step, tr

    print(' Epoch %d @ step %d: Train Loss: %3f, Train Accuracy: %3f' % (
        i, training_step, training_loss, training_acc))
    training_loss, training_acc = 0.0, 0.0

training_step+=1

model.eval()
with torch.no_grad():
    validation_loss, validation_acc = 0.0, 0.0
    count = 0
    for images, labels in test_loader:
        images, labels = images.to(device), labels.to(device)
        output = model(images)
        validation_loss+=loss_func(output,labels)
        validation_acc+=(torch.argmax(output, dim=1) == labels).float().mean()
        count += 1
    validation_loss/=count
    validation_acc/=count

    all_test_loss = np.concatenate((all_test_loss, [[training_step, validation
    all_test_acc = np.concatenate((all_test_acc, [[training_step, validation_a

    epoch_time = time.clock() - epoch_start

    print('Epoch %d Test Loss: %3f, Test Accuracy: %3f, time: %.1fs' % (
        i, validation_loss, validation_acc, epoch_time))

total_time = time.clock() - start
print('Final Test Loss: %3f, Test Accuracy: %3f, Total time: %.1fs' % (
    validation_loss, validation_acc, total_time))

return {'loss': { 'train': all_training_loss, 'test': all_test_loss },
        'accuracy': { 'train': all_training_acc, 'test': all_test_acc }}

def plot_graphs(model_name, metrics):
    for metric, values in metrics.items():
        for name, v in values.items():
            plt.plot(v[:,0], v[:,1], label=name)
        plt.title(f'{metric} for {model_name}')
        plt.legend()
        plt.xlabel("Training Steps")
        plt.ylabel(metric)
        plt.show()

```

Load the **CIFAR-10** dataset and define the transformations. You may also want to print its structure, size, as well as sample a few images to get a sense of how to design the network.

```
!mkdir hw10_data
```

```
mkdir: cannot create directory 'hw10_data': File exists
```

```
# Download the data.
from torchvision import datasets, transforms

transformations = transforms.Compose(
    [transforms.ToTensor(),
     transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))])
train_set = datasets.CIFAR10(root='hw10_data/', download=True, transform=transformations)
test_set = datasets.CIFAR10(root='hw10_data', download=True, train=False, transform=transformations)
```

```
Files already downloaded and verified
Files already downloaded and verified
```

Use `DataLoader` to create a loader for the training set and a loader for the testing set. You can use a `batch_size` of 8 to start, and change it if you wish.

```
from torch.utils.data import DataLoader

batch_size = 256 #8
train_loader = torch.utils.data.DataLoader(train_set, batch_size, shuffle=True, num_workers=4)
test_loader = torch.utils.data.DataLoader(test_set, batch_size, shuffle=True, num_workers=4)

input_shape = np.array(train_set[0][0]).shape
input_dim = input_shape[1]*input_shape[2]*input_shape[0]
```

```
training_epochs = 25
```

▼ Part 1 CIFAR10 with Fully Connected Neural Network (25 Points)

As a warm-up, let's begin by training a two-layer fully connected neural network model on **CIFAR-10** dataset. You may go back to check HW8 for some basics.

We will give you this code to use as a baseline to compare against your CNN models.

```

class TwoLayerModel(nn.Module):
    def __init__(self):
        super(TwoLayerModel, self).__init__()
        self.net = nn.Sequential(
            Flatten(),
            nn.Linear(input_dim, 64),
            nn.ReLU(),
            nn.Linear(64, 10))

    def forward(self, x):
        return self.net(x)

model = TwoLayerModel().to(device)

loss = nn.CrossEntropyLoss()
optimizer = optim.RMSprop(model.parameters(), lr=0.001, weight_decay=0.01)

# Training epoch should be about 15-20 sec each on GPU.
metrics = train(model, train_loader, test_loader, loss, optimizer, training_epoc

```

```



    Epoch 0 @ step 0: Train Loss: 0.004295, Train Accuracy: 0.000113
Epoch 0 Test Loss: 1.799146, Test Accuracy: 0.364551, time: 1.2s
Epoch 1 Test Loss: 1.717270, Test Accuracy: 0.380859, time: 1.3s
Epoch 2 Test Loss: 1.603553, Test Accuracy: 0.427051, time: 1.3s
Epoch 3 Test Loss: 1.658042, Test Accuracy: 0.418457, time: 1.2s
Epoch 4 Test Loss: 1.598114, Test Accuracy: 0.434277, time: 1.2s
    Epoch 5 @ step 1000: Train Loss: 1.620169, Train Accuracy: 0.432607
Epoch 5 Test Loss: 1.552316, Test Accuracy: 0.462598, time: 1.2s
Epoch 6 Test Loss: 1.611885, Test Accuracy: 0.426367, time: 1.2s
Epoch 7 Test Loss: 1.580764, Test Accuracy: 0.444141, time: 1.2s
Epoch 8 Test Loss: 1.697376, Test Accuracy: 0.410938, time: 1.2s
Epoch 9 Test Loss: 1.579937, Test Accuracy: 0.442578, time: 1.2s
    Epoch 10 @ step 2000: Train Loss: 1.526308, Train Accuracy: 0.463935
Epoch 10 Test Loss: 1.559898, Test Accuracy: 0.449414, time: 1.2s
Epoch 11 Test Loss: 1.529849, Test Accuracy: 0.450781, time: 1.3s
Epoch 12 Test Loss: 1.542025, Test Accuracy: 0.448828, time: 1.2s
Epoch 13 Test Loss: 1.528788, Test Accuracy: 0.462012, time: 1.2s
Epoch 14 Test Loss: 1.613499, Test Accuracy: 0.434570, time: 1.2s
    Epoch 15 @ step 3000: Train Loss: 1.502717, Train Accuracy: 0.473353
Epoch 15 Test Loss: 1.578196, Test Accuracy: 0.438281, time: 1.2s
Epoch 16 Test Loss: 1.570323, Test Accuracy: 0.449902, time: 1.2s
Epoch 17 Test Loss: 1.546942, Test Accuracy: 0.450098, time: 1.2s
Epoch 18 Test Loss: 1.586191, Test Accuracy: 0.445312, time: 1.2s
Epoch 19 Test Loss: 1.562466, Test Accuracy: 0.441406, time: 1.2s
    Epoch 20 @ step 4000: Train Loss: 1.493003, Train Accuracy: 0.477532
Epoch 20 Test Loss: 1.529687, Test Accuracy: 0.460449, time: 1.3s
Epoch 21 Test Loss: 1.541251, Test Accuracy: 0.460645, time: 1.2s
Epoch 22 Test Loss: 1.632305, Test Accuracy: 0.427637, time: 1.2s
Epoch 23 Test Loss: 1.567449, Test Accuracy: 0.449707, time: 1.2s
Epoch 24 Test Loss: 1.523372, Test Accuracy: 0.461719, time: 1.2s
Final Test Loss: 1.523372, Test Accuracy: 0.461719, Total time: 30.3s

```

Plot the model results

Normally we would want to use Tensorboard for looking at metrics. However, if colab reset while we are working, we might lose our logs and therefore our metrics. Let's just plot some graphs that will survive across colab instances.

```
plot_graphs("TwoLayerModel", metrics)
```



▼ Part 2 Convolutional Neural Network (CNN) (35 Points)

Now, let's design a convolution neural network!

Build a simple CNN model, inserting 2 CNN layers in from of our 2 layer fully connect model from above:

1. A convolution with 3x3 filter, 16 output channels, stride = 1, padding=1
2. A ReLU activation
3. A Max-Pooling layer with 2x2 window
4. A convolution, 3x3 filter, 16 output channels, stride = 1, padding=1
5. A ReLU activation
6. Flatten layer
7. Fully connected linear layer with output size 64
8. ReLU
9. Fully connected linear layer, with output size 10

You will have to figure out the input sizes of the first fully connected layer based on the previous layer sizes. Note that you also need to fill those in the report section (see report section in the notebook for details)


```

class ConvModel(nn.Module):
    def __init__(self):
        super(ConvModel, self).__init__()
        self.net = nn.Sequential(
            nn.Conv2d(in_channels=3,out_channels=16, kernel_size=3, stride=1,pad
            nn.ReLU(),
            nn.MaxPool2d(kernel_size=2,stride=1),
            nn.Conv2d(in_channels=16,out_channels=16, kernel_size=3, stride=1,pa
            nn.ReLU(),
            Flatten(),
            nn.Linear(31*31*16, 64),
            nn.ReLU(),
            nn.Linear(64, 10),
        )

    def forward(self, x):
        return self.net(x)

model = ConvModel().to(device)

loss = nn.CrossEntropyLoss()
optimizer = optim.RMSprop(model.parameters(), lr=0.001, weight_decay=0.01)

metrics = train(model, train_loader, test_loader, loss, optimizer, training_epoc

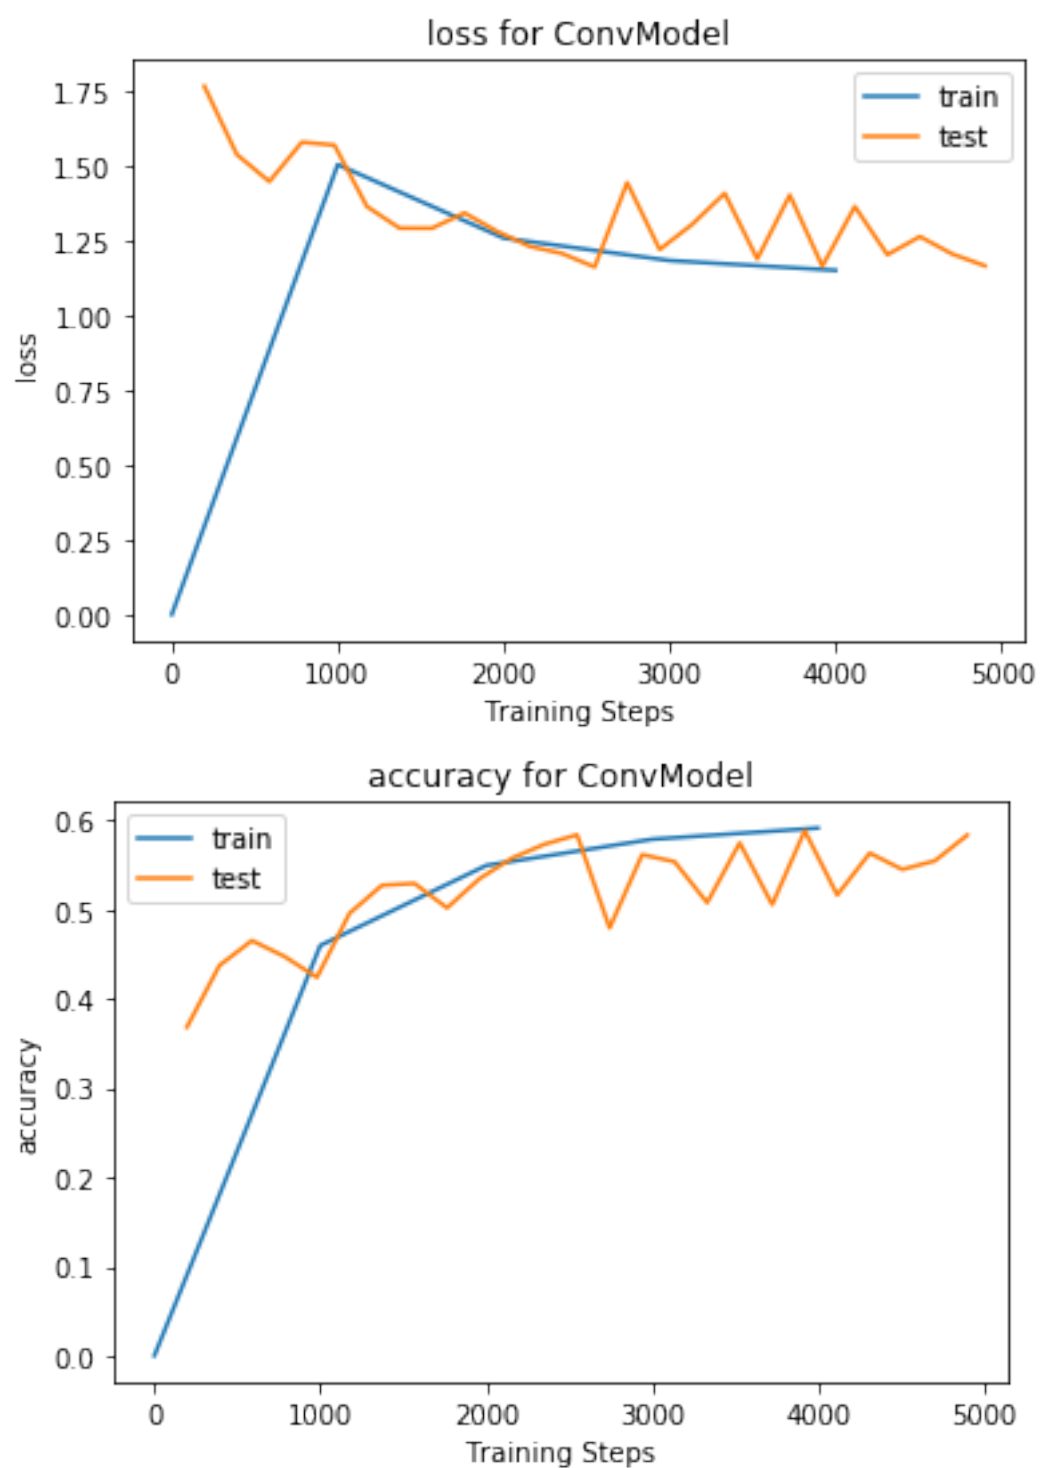
```

```

Epoch 0 @ step 0: Train Loss: 0.004304, Train Accuracy: 0.000102
Epoch 0 Test Loss: 1.766978, Test Accuracy: 0.368066, time: 2.7s
Epoch 1 Test Loss: 1.538963, Test Accuracy: 0.437109, time: 2.7s
Epoch 2 Test Loss: 1.448862, Test Accuracy: 0.465039, time: 3.0s
Epoch 3 Test Loss: 1.580119, Test Accuracy: 0.447266, time: 2.7s
Epoch 4 Test Loss: 1.570375, Test Accuracy: 0.423828, time: 2.9s
Epoch 5 @ step 1000: Train Loss: 1.505343, Train Accuracy: 0.459505
Epoch 5 Test Loss: 1.365523, Test Accuracy: 0.495508, time: 2.8s
Epoch 6 Test Loss: 1.293404, Test Accuracy: 0.526953, time: 2.7s
Epoch 7 Test Loss: 1.293109, Test Accuracy: 0.529004, time: 2.7s
Epoch 8 Test Loss: 1.343846, Test Accuracy: 0.501660, time: 2.7s
Epoch 9 Test Loss: 1.282475, Test Accuracy: 0.534375, time: 2.8s
Epoch 10 @ step 2000: Train Loss: 1.260061, Train Accuracy: 0.549146
Epoch 10 Test Loss: 1.232488, Test Accuracy: 0.557422, time: 2.6s
Epoch 11 Test Loss: 1.208673, Test Accuracy: 0.572949, time: 2.6s
Epoch 12 Test Loss: 1.163653, Test Accuracy: 0.583496, time: 2.7s
Epoch 13 Test Loss: 1.445198, Test Accuracy: 0.479199, time: 2.8s
Epoch 14 Test Loss: 1.221964, Test Accuracy: 0.561328, time: 2.7s
Epoch 15 @ step 3000: Train Loss: 1.185207, Train Accuracy: 0.578513
Epoch 15 Test Loss: 1.306032, Test Accuracy: 0.553320, time: 2.7s
Epoch 16 Test Loss: 1.409430, Test Accuracy: 0.507422, time: 2.7s
Epoch 17 Test Loss: 1.190694, Test Accuracy: 0.574512, time: 2.7s
Epoch 18 Test Loss: 1.403742, Test Accuracy: 0.505371, time: 2.7s
Epoch 19 Test Loss: 1.166340, Test Accuracy: 0.588086, time: 2.6s
Epoch 20 @ step 4000: Train Loss: 1.152617, Train Accuracy: 0.591070
Epoch 20 Test Loss: 1.365484, Test Accuracy: 0.516113, time: 2.8s
Epoch 21 Test Loss: 1.204767, Test Accuracy: 0.562988, time: 2.9s
Epoch 22 Test Loss: 1.265152, Test Accuracy: 0.544727, time: 2.7s
Epoch 23 Test Loss: 1.206752, Test Accuracy: 0.554297, time: 2.7s
Epoch 24 Test Loss: 1.167589, Test Accuracy: 0.583203, time: 2.7s
Final Test Loss: 1.167589, Test Accuracy: 0.583203, Total time: 68.4s

```

```
plot_graphs("ConvModel", metrics)
```



Do you notice the improvement over the accuracy compared to that in Part 1?

▼ Part 3 Open Design Competition (35 Points + 10 bonus points)

Try to beat the previous models by adding additional layers, changing parameters, etc. You should add at least one layer.

Possible changes include:

- Dropout
- Batch Normalization
- More layers
- Residual Connections (harder)
- Change layer size
- Pooling layers, stride
- Different optimizer
- Train for longer

Once you have a model you think is great, evaluate it against our hidden test data (see hidden_loader above) and upload the results to the leader board on gradescope. **The top 3 scorers will get a bonus 10 points.**

You can steal model structures found on the internet if you want. The only constraint is that **you must train the model from scratch.**

```
# You Awesome Super Best model code here
class AwesomeModel(nn.Module):
    def __init__(self):
        super(AwesomeModel, self).__init__()
        self.net = nn.Sequential(
            nn.Conv2d(in_channels=3,out_channels=32, kernel_size=3, stride=1,pad
            nn.ReLU(),
            nn.BatchNorm2d(32),

            nn.Conv2d(in_channels=32,out_channels=32, kernel_size=3, stride=1,pa
            nn.ReLU(),
            nn.BatchNorm2d(32),
            nn.MaxPool2d(kernel_size=2,stride=2),
            nn.Dropout2d(0.2),

            nn.Conv2d(in_channels=32,out_channels=64, kernel_size=3, stride=1,pa
            nn.ReLU(),
            nn.BatchNorm2d(64),
            nn.Conv2d(in_channels=64,out_channels=64, kernel_size=3, stride=1,pa
            nn.ReLU(),
            nn.BatchNorm2d(64),
            nn.MaxPool2d(kernel_size=2,stride=2),
            nn.Dropout2d(0.3),

            nn.Conv2d(in_channels=64,out_channels=128, kernel_size=3, stride=1,p
            nn.ReLU(),
            nn.BatchNorm2d(128),
            nn.Conv2d(in_channels=128,out_channels=128, kernel_size=3, stride=1,
            nn.ReLU(),
            nn.BatchNorm2d(128),
            nn.MaxPool2d(kernel_size=2,stride=2),
            nn.Dropout2d(0.4),

            Flatten(),
            nn.Linear(128*4*4, 10)
        )
```

```
def forward(self, x):  
    return self.net(x)
```

```
model = AwesomeModel().to(device)
```

```
loss = nn.CrossEntropyLoss()
```

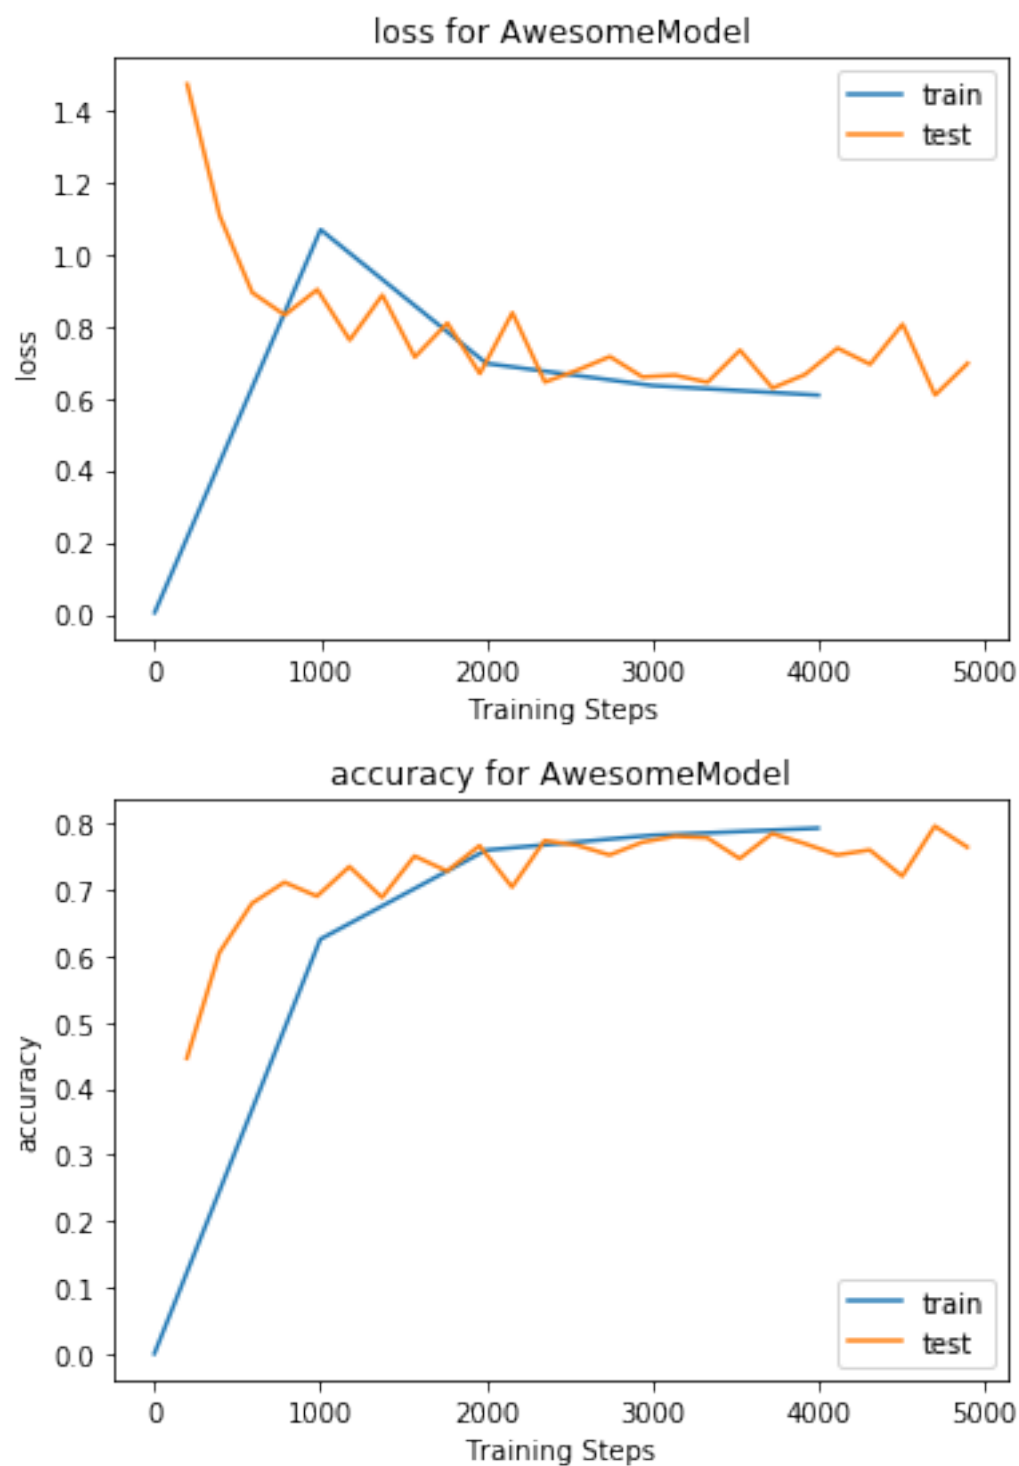
```
optimizer = optim.RMSprop(model.parameters(), lr=0.001, weight_decay=0.01)
```

```
metrics = train(model, train_loader, test_loader, loss, optimizer, training_epoc
```

```
    Epoch 0 @ step 0: Train Loss: 0.004489, Train Accuracy: 0.000094  
Epoch 0 Test Loss: 1.471441, Test Accuracy: 0.445801, time: 6.6s  
Epoch 1 Test Loss: 1.105345, Test Accuracy: 0.605273, time: 6.6s  
Epoch 2 Test Loss: 0.892807, Test Accuracy: 0.679492, time: 6.5s  
Epoch 3 Test Loss: 0.830514, Test Accuracy: 0.711328, time: 6.5s  
Epoch 4 Test Loss: 0.901177, Test Accuracy: 0.690234, time: 6.1s  
    Epoch 5 @ step 1000: Train Loss: 1.067488, Train Accuracy: 0.625320  
Epoch 5 Test Loss: 0.760794, Test Accuracy: 0.734961, time: 6.5s  
Epoch 6 Test Loss: 0.885828, Test Accuracy: 0.688379, time: 6.5s  
Epoch 7 Test Loss: 0.713735, Test Accuracy: 0.750879, time: 6.7s  
Epoch 8 Test Loss: 0.808031, Test Accuracy: 0.728125, time: 6.2s  
Epoch 9 Test Loss: 0.667858, Test Accuracy: 0.766504, time: 6.2s  
    Epoch 10 @ step 2000: Train Loss: 0.696012, Train Accuracy: 0.759795  
Epoch 10 Test Loss: 0.836739, Test Accuracy: 0.704004, time: 6.5s  
Epoch 11 Test Loss: 0.644305, Test Accuracy: 0.774219, time: 6.0s  
Epoch 12 Test Loss: 0.677456, Test Accuracy: 0.767188, time: 6.3s  
Epoch 13 Test Loss: 0.715041, Test Accuracy: 0.752441, time: 6.2s  
Epoch 14 Test Loss: 0.658518, Test Accuracy: 0.771680, time: 6.5s  
    Epoch 15 @ step 3000: Train Loss: 0.634620, Train Accuracy: 0.782181  
Epoch 15 Test Loss: 0.663441, Test Accuracy: 0.780664, time: 6.5s  
Epoch 16 Test Loss: 0.643626, Test Accuracy: 0.778516, time: 6.4s  
Epoch 17 Test Loss: 0.732785, Test Accuracy: 0.746973, time: 6.6s  
Epoch 18 Test Loss: 0.627684, Test Accuracy: 0.784766, time: 6.0s  
Epoch 19 Test Loss: 0.665298, Test Accuracy: 0.769629, time: 6.4s  
    Epoch 20 @ step 4000: Train Loss: 0.607939, Train Accuracy: 0.792910  
Epoch 20 Test Loss: 0.738464, Test Accuracy: 0.752539, time: 6.6s  
Epoch 21 Test Loss: 0.693724, Test Accuracy: 0.760059, time: 6.5s  
Epoch 22 Test Loss: 0.804956, Test Accuracy: 0.720508, time: 6.5s  
Epoch 23 Test Loss: 0.608287, Test Accuracy: 0.796094, time: 6.4s  
Epoch 24 Test Loss: 0.696222, Test Accuracy: 0.764258, time: 6.3s  
Final Test Loss: 0.696222, Test Accuracy: 0.764258, Total time: 159.9s
```

What changes did you make to improve your model?

```
plot_graphs("AwesomeModel", metrics)
```



After you get a nice model, download the test_file.zip and unzip it to get test_file.pt. In colab, you can explore your files from the left side bar. You can also download the files to your machine from there.

```
!wget http://courses.engr.illinois.edu/cs498aml/sp2019/homeworks/test_file.zip
!unzip test_file.zip
```

```
--2019-04-30 15:37:55-- http://courses.engr.illinois.edu/cs498aml/sp2019/
Resolving courses.engr.illinois.edu (courses.engr.illinois.edu)... 130.126
Connecting to courses.engr.illinois.edu (courses.engr.illinois.edu)|130.12
HTTP request sent, awaiting response... 301 Moved Permanently
Location: https://courses.engr.illinois.edu/cs498aml/sp2019/homeworks/test
--2019-04-30 15:37:55-- https://courses.engr.illinois.edu/cs498aml/sp2019
Connecting to courses.engr.illinois.edu (courses.engr.illinois.edu)|130.12
HTTP request sent, awaiting response... 200 OK
Length: 3841776 (3.7M) [application/x-zip-compressed]
Saving to: 'test_file.zip.1'
```

```
test_file.zip.1      100%[=====>]    3.66M  14.0MB/s    in 0.3
```

```
2019-04-30 15:37:55 (14.0 MB/s) - 'test_file.zip.1' saved [3841776/3841776]
```

```
Archive: test_file.zip
replace test_file.pt? [y]es, [n]o, [A]ll, [N]one, [r]ename: y
  inflating: test_file.pt
```

Then use your model to predict the label of the test images. Fill the remaining code below, where `x` has two dimensions (batch_size x one image size). Remember to reshape `x` accordingly before feeding it into your model. The submission.txt should contain one predicted label (0~9) each line. Submit your submission.txt to the competition in gradscope.

```
import torch.utils.data as Data

test_file = 'test_file.pt'
pred_file = 'submission.txt'
f_pred = open(pred_file, 'w')
tensor = torch.load(test_file)
torch_dataset = Data.TensorDataset(tensor)
test_loader = torch.utils.data.DataLoader(torch_dataset, batch_size, shuffle=False)
for ele in test_loader:
    x = ele[0]
    x = x.to(device)
    x = x.reshape(x.shape[0], 3, 32, 32)
    _, pred = torch.max(model(x), 1)
    pred = pred.cpu().detach().numpy()
    for i in enumerate(pred):
        f_pred.write(str(i[1]))
        f_pred.write('\n')
f_pred.close()
```

Report

Part 0: Imports and Basic Setup (5 Points)

Nothing to report for this part. You will be just scored for finishing the setup.

Part 1: Fully connected neural networks (25 Points)

Test (on validation set) accuracy (5 Points):0.461719


Test loss (5 Points):1.523372

Training time (5 Points):30.3s


Plots:

- Plot a graph of accuracy on validation set vs training steps (5 Points) - **Shown Below**
- Plot a graph of loss on validation set vs training steps (5 Points) - **Shown Below**

```
#Plot a graph of accuracy on validation set vs training steps (5 Points)
from IPython.display import Image
Image(filename="2layer_accuracy.png",width=600,height=400)
```




```
#Plot a graph of accuracy on validation set vs training steps (5 Points)
from IPython.display import Image
Image(filename="2layer_loss.png",width=600,height=400).
```



▼ Part 2: Convolution Network (Basic) (35 Points)

Tensor dimensions: A good way to debug your network for size mismatches is to print the dimension of output after every layers:

(10 Points)

Output dimension after 1st conv layer: [256, 16, 32, 32]

Output dimension after 1st max pooling: [256, 16, 31, 31]

Output dimension after 2nd conv layer: [256, 16, 31, 31]

Output dimension after flatten layer: [256, 15376]

Output dimension after 1st fully connected layer: [256, 64]

Output dimension after 2nd fully connected layer: [256, 10]

Test (on validation set) Accuracy (5 Points):0.583203

Test loss (5 Points):1.167589

Training time (5 Points):68.4s


Plots:

- Plot a graph of accuracy on validation set vs training steps (5 Points) - **Shown Below**
- Plot a graph of loss on validation set vs training steps (5 Points) - **Shown Below**

#Plot a graph of accuracy on validation set vs training steps (5 Points)

```
from IPython.display import Image
```


```
Image(filename="conv_accuracy.png",width=600,height=400)
```



#Plot a graph of accuracy on validation set vs training steps (5 Points)

```
from IPython.display import Image
```

```
Image(filename="conv_loss.png",width=600,height=400.)
```



Part 3: Convolution Network (Add one or more suggested changes) (35 Points)

Describe the additional changes implemented, your intuition for as to why it works, you may also describe other approaches you experimented with (10 Points):

Below are the addition:

- (1) Batch Normalization:** This layer doesn't have any activation function, hence will have higher learning rate, so it will be able to detect the pattern (via training), where the previous layer was unable to learn it.
- (2) Drop Out Layer:** This will decrease the overfitting of the training data. Every Layer I am dropping 10% additional neuron, which also in turn removes few information from the layer. However it will benefit to remove few neurons which have influenced on the training to become overfitting.
- (3) Pooling Layer:** This will reduce the size of the dimension to half ($1/2$) after each pooling layer, which reduces the overall time to train the model and avoid the curse of dimensionality.
- (4) More Layer to the Network:** More Layer being added up, hence each layer will understand/generate at least some part of the feature (image). This will increase the accuracy of the test. Though it will not always be true to increase the accuracy by adding more layer (because each layer is not learning anything new feature), however in my Awesome model, will have 6 Layer which produce the accuracy of 78% in the Test Set.
- (5) Increase the Batch Size:** Batch_size is set to 256, so for one pass to the network, it will take 256 random samples, which will increase the estimate of the gradient for each forward pass.
- (6) Increase the number of epoch:** Since the parameters get updated after each forward pass per batch_size (in my case it's 256) till last samples are being passed, and it will get repeated till all epochs. The larger the epoch, the longer it will train and it will reach to the global minima for more number of epoch, as the parameters get updated on each forward pass and per each epoch. However this will not be always true, as it could oscillate around the global minima after it reaches to certain epoch and never converges to global minima. I have chosen the number of epoch as 25.

Test (on validation set) Accuracy (5 Points):0.764258

Test loss (5 Points):0.696222


Training time (5 Points):159.9s

Plots:

- Plot a graph of accuracy on validation set vs training steps (5 Points) - **Shown Below**
- Plot a graph of loss on validation set vs training steps (5 Points) - **Shown Below**

10 bonus points will be awarded to top 3 scorers on leaderboard (in case of tie for 3rd position everyone tied for 3rd position will get the bonus)

```
#Plot a graph of accuracy on validation set vs training steps (5 Points)
from IPython.display import Image
Image(filename="awesome_accuracy.png",width=600,height=400.)
```



```
#Plot a graph of accuracy on validation set vs training steps (5 Points)
from IPython.display import Image
Image(filename="awesome_loss.png",width=600,height=400.)
```

